

# BIOSTAT C161 HW3

:)

Hanxi Chen

2025-10-31

```
library(tidyverse)
```

Warning: package 'ggplot2' was built under R version 4.4.3

```
library(dplyr)
library(caret)
```

## Q1

(iii)

Plot the function  $f$  over  $x \in [-2.5, 2.5]$ . How many local/global minima do you see? What are their approximate values? Can there be other local minima?

**Sol:**

```
# Clear environment
rm(list = ls())

f <- function(x) {
  x^4 - 6*x^2 + 4*x + 18
}

x_vals <- seq(-2.5, 2.5, by = 0.01)
y_vals <- f(x_vals)

plot(x_vals, y_vals, type = "l", col = "blue", lwd = 2,
     main = expression(f(x) == x^4 - 6*x^2 + 4*x + 18),
     xlab = "x", ylab = "f(x)")

grid()
f_prime <- function(x) 4*x^3 - 12*x + 4
roots <- uniroot.all <- function(f, interval, n = 1000, ...) {
```

```


xseq <- seq(interval[1], interval[2], len = n)
yseq <- f(xseq)
sign_change <- which(diff(sign(yseq)) != 0)
sapply(sign_change, function(i)
  uniroot(f, c(xseq[i], xseq[i + 1]), ...)$root)
}

critical_points <- uniroot.all(f_prime, c(-2.5, 2.5))
critical_points

[1] -1.8793794  0.3472973  1.5320821

points(critical_points, f(critical_points), col = "red", pch = 19)
text(critical_points, f(critical_points) + 2,
     labels = round(critical_points, 2), col = "red")

```



```

data.frame(
  x = critical_points,
  f_x = f(critical_points)
)

```

	x	f_x
1	-1.8793794	1.765578
2	0.3472973	18.680045
3	1.5320821	15.554378

The plot of the function over the interval  $[-2.5, 2.5]$  shows three stationary

points at approximately  $x = -1.88$ ,  $0.35$ , and  $1.53$ . Based on their function values  $f(-1.88) \approx 1.77$ ,  $f(0.35) \approx 18.68$ , and  $f(1.53) \approx 15.55$ , the points at  $x = -1.88$  and  $x = 1.53$  are local minima, while  $x = 0.35$  is a local maximum. The global minimum occurs at  $x = -1.88$  because it gives the smallest  $f(x)$ . There cannot be any additional local minima since the derivative  $4x^3 - 12x + 4$  is cubic and thus has at most three real roots.

(iv)

Write your own code for  $S$  steps of GD on this function (do not use built-in or third-party GD codes).

**Sol:**

```
# Gradient Descent
gradient_descent <- function(x0, alpha, S) {
  x_vals <- numeric(S + 1) # store all iterates
  x_vals[1] <- x0           # initial value
  grad_vals <- numeric(S + 1) # store gradients

  for (s in 1:S) {
    grad <- f_prime(x_vals[s])
    grad_vals[s] <- grad
    x_vals[s + 1] <- x_vals[s] - alpha * grad
  }

  grad_vals[S + 1] <- f_prime(x_vals[S + 1]) # gradient at last step

  return(data.frame(Step = 0:S,
                    x = x_vals,
                    f_x = f(x_vals),
                    gradient = grad_vals))
}
```

(v)

Repeat (i) and (ii) using your code with  $S = 20$  steps. Do you observe convergence in both cases?

**Sol:**

```
# Case 1: Start from  $x(0) = 1$ 
result_1 <- gradient_descent(x0 = 1, alpha = 0.1, S = 20)
print(result_1)
```

	Step	x	f_x	gradient
1	0	1.000000	17.000000	-4.000000000000

```

2    1 1.400000 15.68160 -1.8240000000
3    2 1.582400 15.57563  0.8604535849
4    3 1.496355 15.56442 -0.5544413756
5    4 1.551799 15.55757  0.3258336573
6    5 1.519215 15.55570 -0.2050942422
7    6 1.539725 15.55485  0.1245284454
8    7 1.527272 15.55456 -0.0774512605
9    8 1.535017 15.55445  0.0475001130
10   9 1.530267 15.55440 -0.0293927124
11  10 1.533206 15.55439  0.0180900918
12  11 1.531397 15.55438 -0.0111713403
13  12 1.532515 15.55438  0.0068845274
14  13 1.531826 15.55438 -0.0042481318
15  14 1.532251 15.55438  0.0026192715
16  15 1.531989 15.55438 -0.0016157492
17  16 1.532151 15.55438  0.0009964086
18  17 1.532051 15.55438 -0.0006145838
19  18 1.532112 15.55438  0.0003790316
20  19 1.532074 15.55438 -0.0002337761
21  20 1.532098 15.55438  0.0001441804

```

```

# Case 2: Start from x(0) = 0
result_2 <- gradient_descent(x0 = 0, alpha = 0.1, S = 20)
print(result_2)


```

	Step	x	f_x	gradient
1	0	0.0000000	18.000000	4.000000
2	1	-0.4000000	15.465600	8.544000
3	2	-1.2544000	6.017247	11.157509
4	3	-2.3701509	6.371307	-20.816574
5	4	-0.2884935	16.353582	7.365878
6	5	-1.0250813	8.699088	11.992388
7	6	-2.2243201	3.895865	-13.328344
8	7	-0.8914858	10.297198	11.863807
9	8	-2.0778665	2.424417	-6.950597
10	9	-1.3828067	4.652181	10.017121
11	10	-2.3845188	6.676144	-21.618601
12	11	-0.2226587	16.814362	6.627750
13	12	-0.8854337	10.368953	11.848510
14	13	-2.0702846	2.372859	-6.650193
15	14	-1.4052653	4.430034	9.762877
16	15	-2.3815530	6.612276	-21.452085
17	16	-0.2363445	16.722590	6.783327
18	17	-0.9146772	10.021442	11.915125
19	18	-2.1061897	2.637438	-8.098250
20	19	-1.2963648	5.555459	10.841894

```
21    20 -2.3805541  6.590875 -21.396114
```

```
# Plot convergence paths
plot(result_1$Step, result_1$x, type = "l", col = "blue", lwd = 2,
     ylim = range(c(result_1$x, result_2$x)),
     main = "Convergence Paths of Gradient Descent (  $\alpha = 0.1$ ,  $S = 20$ )",
     xlab = "Step", ylab = "x")
lines(result_2$Step, result_2$x, col = "red", lwd = 2)
legend("topright", legend = c("x(0)=1", "x(0)=0"), col = c("blue", "red"), lwd = 2)
```

### Convergence Paths of Gradient Descent ( $\alpha = 0.1$ , $S = 20$ )



The plot shows the convergence behavior of gradient descent with  $\alpha = 0.1$  and  $S = 20$  for two different initial points. When starting from  $x(0) = 1$  (blue line), the algorithm quickly stabilizes near  $x = 1.53$ , and the final gradient value of 0.000144 indicates that it has effectively converged to a local minimum. In contrast, when starting from  $x(0) = 0$  (red line), the sequence oscillates strongly between negative and positive values, showing that the updates overshoot due to a large learning rate. The final gradient of  $-21.396$  further confirms divergence rather than convergence. Therefore, gradient descent converges successfully only for the initial value  $x(0) = 1$ , while it fails to converge when starting from  $x(0) = 0$  under the same learning rate.

(vi)

Now set the learning rate to  $\alpha = 0.01$  and repeat (v). Explain why GD performs differently from (v)

Sol:

```
# Case 1: Starting from  $x(0) = 1$ 
result_x1_small_alpha <- gradient_descent(x0 = 1, alpha = 0.01, S = 20)
print(result_x1_small_alpha)
```

	Step	x	f_x	gradient
1	0	1.000000	17.00000	-4.000000
2	1	1.040000	16.84026	-3.980544
3	2	1.079805	16.68285	-3.921540
4	3	1.119021	16.53086	-3.823264
5	4	1.157253	16.38715	-3.687701
6	5	1.194130	16.25416	-3.518495
7	6	1.229315	16.13374	-3.320738
8	7	1.262523	16.02703	-3.100610
9	8	1.293529	15.93447	-2.864928
10	9	1.322178	15.85583	-2.620647
11	10	1.348385	15.79033	-2.374400
12	11	1.372129	15.73679	-2.132113
13	12	1.393450	15.69379	-1.898738
14	13	1.412437	15.65982	-1.678116
15	14	1.429218	15.63336	-1.472961
16	15	1.443948	15.61304	-1.284931
17	16	1.456797	15.59762	-1.114766
18	17	1.467945	15.58604	-0.962461
19	18	1.477570	15.57742	-0.827444
20	19	1.485844	15.57106	-0.708743
21	20	1.492931	15.56641	-0.605129

```
# Case 2: Starting from  $x(0) = 0$ 
result_x0_small_alpha <- gradient_descent(x0 = 0, alpha = 0.01, S = 20)
print(result_x0_small_alpha)
```

	Step	x	f_x	gradient
1	0	0.000000	18.00000	4.000000
2	1	-0.040000	17.83040	4.479744
3	2	-0.084797	17.61771	5.015130
4	3	-0.134948	17.35127	5.609555
5	4	-0.191044	17.01816	6.264641
6	5	-0.253690	16.60322	6.978979
7	6	-0.323480	16.08919	7.746370
8	7	-0.400944	15.45752	8.553513
9	8	-0.486479	14.69011	9.377227
10	9	-0.580251	13.77220	10.181555
11	10	-0.682067	12.69686	10.915573
12	11	-0.791222	11.47082	11.513346
13	12	-0.906356	10.12051	11.898055
14	13	-1.025336	8.69602	11.992231

```

15 14 -1.14525919 7.269595 11.734537
16 15 -1.26260457 5.925939 11.100028
17 16 -1.37360485 4.744817 10.116441
18 17 -1.47476926 3.781640 8.867067
19 18 -1.56343992 3.054993 7.474937
20 19 -1.63818929 2.547311 6.072872
21 20 -1.69891801 2.217251 4.772516


```

```

plot(result_x1_small_alpha$Step, result_x1_small_alpha$x,
     type = "b", col = "blue", pch = 19,
     ylim = range(c(result_x1_small_alpha$x, result_x0_small_alpha$x)),
     main = "Convergence Paths of Gradient Descent (  $\alpha$  = 0.01, S = 20)",
     xlab = "Step", ylab = "x")
lines(result_x0_small_alpha$Step, result_x0_small_alpha$x,
     type = "b", col = "red", pch = 17)
legend("bottomright", legend = c("x(0)=1", "x(0)=0"),
     col = c("blue", "red"), pch = c(19, 17))

```

### Convergence Paths of Gradient Descent ( $\alpha = 0.01$ , S = 20)



With a smaller learning rate ( $\alpha = 0.01$ ), the updates become smoother and more stable. The path starting from  $x(0) = 1$  moves slowly toward the local minimum near  $x = 1.53$ , while the path from  $x(0) = 0$  moves gradually toward the global minimum near  $x = -1.88$ . However, after 20 steps, neither starting point has reached a point where the gradient is close to 0, indicating that convergence is not yet achieved due to the slower update rate.

## Q2

The College.csv dataset contains admissions data for a sample of 777 universities. We want to predict the number of applications received (“Apps”) using the other variables in the dataset

(i)

Let the first 600 observations be the training set and the remaining 177 observations be the test set

Sol:

```
college <- read.csv("./College.csv")
```

```
head(college)
```

		X	Private	Apps	Accept	Enroll	Top10perc	Top25perc
1	Abilene Christian University	Yes	1660	1232	721	23	52	
2	Adelphi University	Yes	2186	1924	512	16	29	
3	Adrian College	Yes	1428	1097	336	22	50	
4	Agnes Scott College	Yes	417	349	137	60	89	
5	Alaska Pacific University	Yes	193	146	55	16	44	
6	Albertson College	Yes	587	479	158	38	62	

	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	Personal	PhD	Terminal
1	2885	537	7440	3300	450	2200	70	78
2	2683	1227	12280	6450	750	1500	29	30
3	1036	99	11250	3750	400	1165	53	66
4	510	63	12960	5450	450	875	92	97
5	249	869	7560	4120	800	1500	76	72
6	678	41	13500	3335	500	675	67	73

	S.F.Ratio	perc.alumni	Expend	Grad.Rate
1	18.1	12	7041	60
2	12.2	16	10527	56
3	12.9	30	8735	54
4	7.7	37	19016	59
5	11.9	2	10922	15
6	9.4	11	9727	55

```
college <- college[, -1]
college$Private <- as.factor(college$Private)
train <- college[1:600, ]
test <- college[601:777, ]

dim(train)
```

```
[1] 600 18
```



```
dim(test)
```

```
[1] 177 18
```

```
train_x <- model.matrix(Apps ~ ., data = train)[, -1] # remove intercept
test_x <- model.matrix(Apps ~ ., data = test)[, -1]
train_y <- train$Apps
test_y <- test$Apps
```

(ii)

fit the OLS regression on the training set, and report the test error obtained. For the rest of the problem, let the penalization parameter vary on the 1000-point grid from 0.01 to 60 .

**Sol:**

```
library(glmnet)

# Fit OLS model (lambda = 0 no regularization)
ols_fit <- lm(Apps ~ ., data = train)

ols_pred <- predict(ols_fit, newdata = test)

ols_mse <- mean((ols_pred - test_y)^2)
ols_mse
```

```
[1] 1502077
```

```
sqrt(ols_mse)
```

```
[1] 1225.593
```

```
lambda_grid <- seq(0.01, 60, length = 1000)
```

The test MSE = 1502077, RMSE is 1225.593.

(iii)

Fit the LASSO regression on the training set, with the penalization parameter chosen by 20-fold cross-validation. Report the test error obtained

**Sol:**

```
# alpha = 1 for LASSO
set.seed(123)
lasso_cv <- cv.glmnet(train_x, train_y, alpha = 1, lambda = lambda_grid, nfolds = 20)

# Best lambda
```

```
best_lambda_lasso <- lasso_cv$lambda.min
best_lambda_lasso
```

```
[1] 0.01
```

```
lasso_pred <- predict(lasso_cv, s = best_lambda_lasso, newx = test_x)
lasso_mse <- mean((lasso_pred - test_y)^2)
lasso_rmse <- sqrt(lasso_mse)
```

```
lasso_mse
```

```
[1] 1499849
```

```
lasso_rmse
```

```
[1] 1224.683
```

The optimal penalty parameter selected by 20-fold cross-validation is  $\lambda = 0.01$ , yielding a test MSE of approximately 1,499,849 and a corresponding RMSE of about 1224.68, indicating the average prediction error in the number of applications is around 1225.

(iv)

Fit the ridge regression on the training set, with the penalization parameter chosen by leave-one-out cross-validation. Report the test error obtained. **Sol:**

```
set.seed(123)
```

```
ridge_cv <- cv.glmnet(train_x, train_y, alpha = 0, lambda = lambda_grid,
                      nfolds = nrow(train))
```

Warning: Option grouped=FALSE enforced in cv.glmnet, since < 3 observations per fold

```
# Best lambda
best_lambda_ridge <- ridge_cv$lambda.min
best_lambda_ridge
```

```
[1] 0.01
```

```
ridge_pred <- predict(ridge_cv, s = best_lambda_ridge, newx = test_x)
ridge_mse <- mean((ridge_pred - test_y)^2)
ridge_rmse <- sqrt(ridge_mse)
```

```
ridge_mse
```

```
[1] 1501329
```

```
ridge_rmse
```

[1] 1225.287

The optimal  $\lambda$  selected by leave-one-out cross-validation is 0.01, giving a test MSE of approximately 1,501,329 and an RMSE of about 1225.29. This test error is nearly identical to that of the LASSO model, suggesting that both methods achieve similar predictive performance on this dataset.

(v)

which of the three models do you prefer? Is there much difference among the test errors?

**Sol:** All three models: OLS, LASSO, and ridge regression have very similar test errors, with RMSE values all around 1225. This indicates that neither penalization method provides a meaningful improvement over ordinary least squares for predicting the number of applications. Since the predictive performance is nearly identical, the OLS model would be preferred for its simplicity and ease of interpretation, although LASSO could still be useful for variable selection if model sparsity is desired.